



# Concurrency in databases

Alexander Shelemin

Senior developer at Veeam

# Veeam In Numbers

## R&D Offices :

Prague

Saint - Petersburg



**82%**  
of Fortune 500



**35+**  
Countries where Veeam  
has offices



**400,000+**  
Customers worldwide



**4500+**  
Employees worldwide



**В 3.5 раза**  
Industry Average  
for Customer Satisfaction



**\$1+ млрд.**  
in revenue in 2019

Products Employees



# Concurrency in databases


Concurrency: 2 or more clients executing code at the same time.

Databases: the talk is relevant for popular RDBMS (SQL Server, Postgres, MySQL, Oracle)

# Agenda

- Transactions and atomicity
- Locking and deadlocks
- Isolation levels
- Let's race!

# Transactions



# Transactions: atomicity

atomicity in java, c#, c++: 'no race conditions'

atomicity in databases: 'all or nothing'

Don't mix these!

# Transactions: atomicity

Advice that doesn't really work:

“Just wrap it all in a transaction!”

“Just rewrite everything as a single SQL query”

# Locks, blocking, deadlocks

## Locks:

- type
  - shared
  - exclusive
- scope
  - table
  - page
  - row
- duration



# Locks, blocking, deadlocks

Blocking: progress can be made.

Deadlocks: no progress can be made. DB will abort one of the transactions.

- usually OK to retry

# Transactions: isolation



## Expectation:

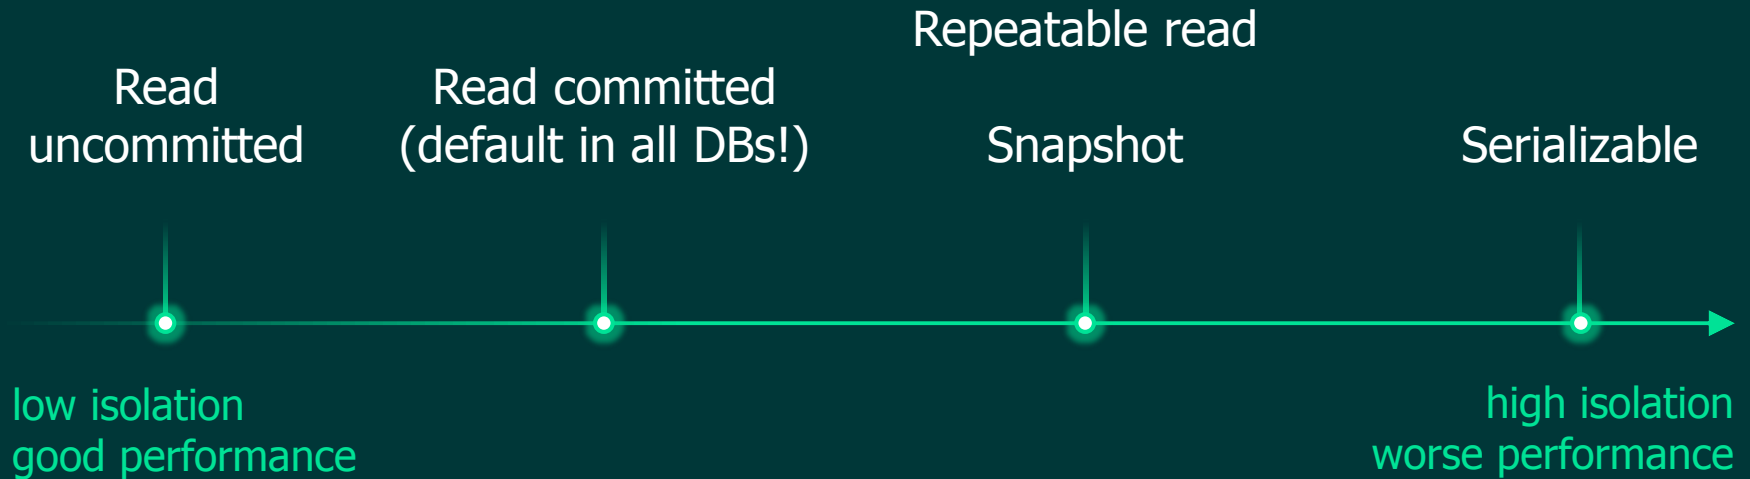
each transaction is  
running as if alone.



## Reality:

performance sucks,  
need some tradeoffs.

# Transactions: isolation levels



# Isolation levels - anomalies

	Possible phenomena (as defined in ANSI SQL Standard)		
Isolation level	Dirty read	Nonrepeatable read	Phantom
Read uncommitted	Yes	Yes	Yes
Read committed	No	Yes	Yes
Read committed using row versioning	No	Yes	Yes
Repeatable read	No	No	Yes
Snapshot	No	No	No
Serializable	No	No	No

Картинка взята из статьи <https://sqlperformance.com/2014/06/sql-performance/the-snapshot-isolation-level>

# Transactions: isolation levels

- Part of SQL language standard (SQL-92)
- Can be set on a transaction level.
- Some systems allow changing the default on DB level, some don't

# Isolation levels: implementation

SQL standard doesn't specify implementation details, so different vendors do things differently.

Two popular approaches:

Locking (read/write locks, 2PL).

- Amount and duration of locks depend on isolation level
- Usually means writers block readers.

MVCC.

- Readers select data from a snapshot.
- Data may be stale.

# Isolation levels: read uncommitted

- Allows dirty reads (i.e. reads of data from transactions that are not yet committed)
- 'NOLOCK' hint in SQL Server – not needed with Read Committed Snapshot.
- In Postgres and Oracle it maps to read committed – no way to read dirty data.
- Just don't use it!

# Isolation levels: read committed

Default in most popular DBMS

## Allows nonrepeatable (fuzzy) reads

- You do the same SELECT twice in the same transaction.
- Rows can be UPDATED or DELETED by other transactions between your SELECTs.

```
-- Session 1
BEGIN TRAN
SELECT num from T1; -- 1

SELECT num from T1; -- 123
COMMIT;
```

```
-- Session 2

UPDATE T1 set num = 123;
```



# Isolation levels: read committed

Default in most popular DBMS (Postgres, SQL Server, Oracle, MySQL)

## 2 different implementations in SQL Server:

- Read Committed – uses locking
- Read Committed Snapshot – uses MVCC.
  - Use Read Committed Snapshot as default!

# Isolation levels: repeatable read

Data that you read cannot change until commit.

But, allows 'phantom reads'

- You do the same SELECT with twice in the same transaction.
- Rows you saw once will stay the same.
- BUT new rows can be INSERTED by other transactions between your SELECTS.

# Isolation levels: repeatable read

```
-- Session 1  
BEGIN TRAN  
SELECT count(*) from T1 where id > 100; -- 666  
  
SELECT count(*) from T1 where id > 100; -- 667  
COMMIT;
```

```
-- Session 2  
  
INSERT INTO T1(id) values (100500);
```

# Isolation levels: snapshot

Provides point-in-time snapshot for the whole database for the duration of transaction.

- Very convenient for reports or batch selects.
- Data can be out of date.
- Allows “write skew” anomaly.
- Read committed snapshot works per statement!

# Isolation levels: serializable

Doesn't allow any anomalies.

Causes a lot of deadlocks/conflicts to guarantee that.

Use sparingly in highly concurrent systems.

# Beyond isolation levels

“SELECT FOR UPDATE” (aka UPDLOCK in SQL Server)

- Useful for concurrent select/update queries.

You can lock the whole table – doesn’t scale, but in rare cases could be a good solution – think batch inserts.

ON CONFLICT (Postgres, MySQL)

Application locks (called advisory locks in Postgres)



Let's UPSERT some data!



Demo



# Conclusions

- Don't panic, but understand that race condition can happen.
- Look for cases with highly concurrent data modifications and pick isolation level and/or locks accordingly.
- Constraints are important! UNIQUE will help you find a problem in testing, not in production when you have duplicate rows.
- Stress testing combined with constraints is powerful.

# References

- Martin Kleppmann's talk on transactions <https://www.youtube.com/watch?v=5ZjhNTM8XU8>, also his book is superb ("Designing data intensive applications").
- Critique of isolation levels paper overview: <https://blog.acolyer.org/2016/02/24/a-critique-of-ansi-sql-isolation-levels/>
- <https://github.com/ept/hermitage> a tool to research transactional anomalies and guarantees.
- <https://sqlperformance.com/2014/07/t-sql-queries/isolation-levels> in-depth explanation of SQL Server isolation levels behavior.
- <http://www.bailis.org/blog/understanding-weak-isolation-is-a-serious-problem/> - interesting thoughts on prevalence of weak isolation in production systems.



Senior C# Developer  
(с возможностью  
переезда в Прагу)

2



Full Stack C# Developer



C# Developer (с  
возможностью релокации в  
Прагу)



Thank you

veeAM