



Обзор перфомансных изменений в .NET 5

Рома Неволин

DevRel в Контуре

Launch of .NET 5

.NET 5 performance improvements

TechEmpower benchmarks round 20 predicted perf increases for .NET 5 over .NET Core 3.1

+38%

11,712,405 RPS

Plaintext

+42%

1,243,436 RPS

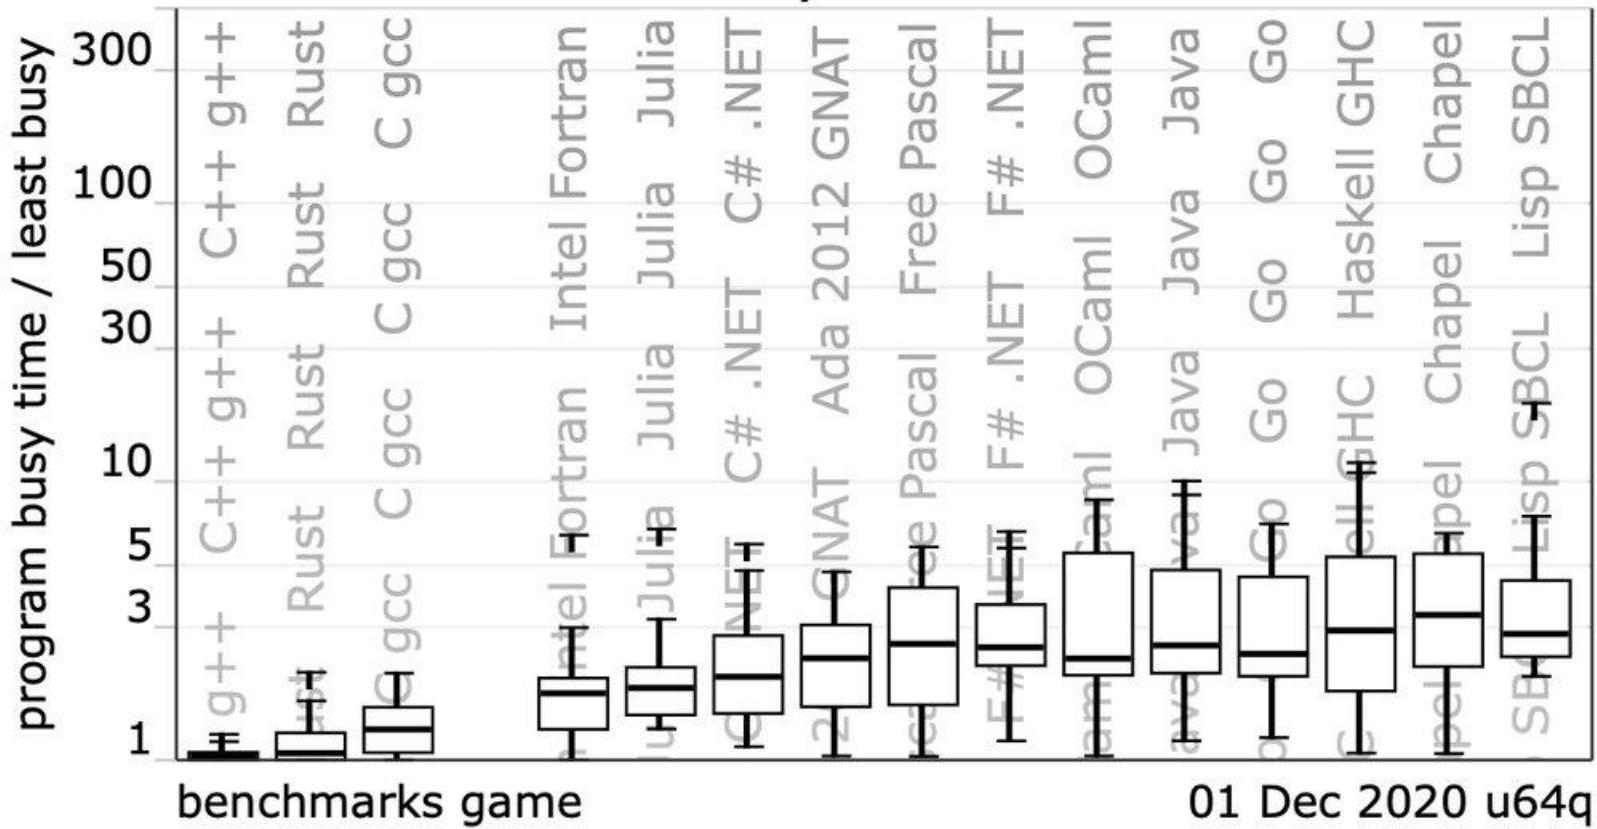
Json

+20%

417,818 RPS

Fortunes

How many times slower?



Performance Improvements in .NET 5



Stephen

July 13th, 2020

In previous releases of .NET Core, I've blogged about the significant performance improvements that found their way into the release. For each post, from [.NET Core 2.0](#) to [.NET Core 2.1](#) to [.NET Core 3.0](#), I found myself having more and more to talk about. Yet interestingly, after each I also found myself wondering whether there'd be enough meaningful improvements next time to warrant another post. Now that .NET 5 is shipping preview releases, I can definitively say the answer is, again, "yes". .NET 5 has already seen a wealth of performance improvements, and even though it's not scheduled for final release until [later this year](#) and there's very likely to be a lot more improvements that find their way in by then, I wanted to highlight a bunch of the improvements that are already available now. In this post, I'll highlight ~250 pull requests that have contributed to myriad of performance improvements across .NET 5.

Работа с текстом

[Benchmark]

```
public int Trim() => " test ".AsSpan().Trim().Length;
```

Method	Runtime	Mean	Ratio	Code Size
Trim	.NET FW 4.8	21.694 ns	1.00	569 B
Trim	.NET Core 3.1	8.079 ns	0.37	377 B
Trim	.NET 5.0	6.556 ns	0.30	365 B

Perf improvements to char.IsWhiteSpace / char.ToUpper / char.ToLower #26848

Open with ▾

Merged

GrabYourPitchfo... merged 3 commits into `dotnet:master` from `GrabYourPitchforks:char_updates` on 24 Sep 2019

Conversation 11

Commits 3

Checks 55

Files changed 1

+51 -86



GrabYourPitchforks commented on 24 Sep 2019

Member



...

This is especially noticeable in benchmarks for APIs like `ReadOnlySpan.Trim()`.

Method	Toolchain	input	Mean	Error	StdDev	Median	Ratio	RatioSD
Trim	master	{empty}	3.509 ns	0.0475 ns	0.0444 ns	3.501 ns	1.00	0.00
Trim	proto	{empty}	3.423 ns	0.0801 ns	0.0710 ns	3.399 ns	0.98	0.03
Trim	master	{spaces on both sides}	9.161 ns	0.2052 ns	0.2195 ns	9.091 ns	1.00	0.00
Trim	proto	{spaces on both sides}	7.299 ns	0.1714 ns	0.1519 ns	7.287 ns	0.79	0.02

Reviewers

- marek-safar
- gfoidl
- adamsitnik
- tarekgh



Assignees

No one assigned

Labels

None yet

Projects

None yet

Improve char.ToUpperInvariant / ToLowerInvariant perf #35194

Open with ▾

Merged

GrabYourPitchfo... merged 1 commit into [dotnet:master](#) from [GrabYourPitchforks:invariant_perf](#) on 20 Apr 2020

Conversation 12

Commits 1

Checks 125

Files changed 2

+32 -17



GrabYourPitchforks commented on 20 Apr 2020

Member ...

Somewhat inspired by Steve's PR at [#35185](#). Turns out we call `char.ToUpperInvariant` / `char.ToLowerInvariant` in a few dozen places within the runtime.

This PR improves these code paths by: (a) not bouncing through the `TextInfo.Invariant` indirection for ASCII data or when `GlobalizationMode.Invariant` is enabled; (b) inlining small methods that the JIT wasn't already inlining; and (c) tweaking the bit twiddling to get more efficient codegen.

Method	Toolchain	args	Mean	Error	StdDev	Ratio	Code Size
CharToUpperInvariant	master	Hello World!	51.98 ns	0.615 ns	0.545 ns	1.00	147 B
CharToUpperInvariant	textinfo	Hello World!	22.71 ns	0.436 ns	0.408 ns	0.44	105 B
CharToLowerInvariant	master	Hello World!	49.15 ns	0.295 ns	0.247 ns	1.00	144 B
CharToLowerInvariant	textinfo	Hello World!	22.68 ns	0.242 ns	0.215 ns	0.46	105 B

Reviewers

- gfoidl
- stephentoub
- tarekgh



Assignees

No one assigned

Labels

[area-System.Runtime](#) [tenet-performance](#)

Projects

None yet

[Benchmark]

```
[Arguments("It's exciting to see great performance!")]
public int ToUpperInvariant(string s)
{
    int sum = 0;

    for (int i = 0; i < s.Length; i++)
        sum += char.ToUpperInvariant(s[i]);

    return sum;
}
```

Method	Runtime	Mean	Ratio	Code Size
ToUpperInvariant	.NET FW 4.8	208.34 ns	1.00	171 B
ToUpperInvariant	.NET Core 3.1	166.10 ns	0.80	164 B
ToUpperInvariant	.NET 5.0	69.15 ns	0.33	105 B

Optimize integral ToString (without generics) #32528

Merged

jkotas merged 1 commit into `dotnet:master` from `ts2do:IntFormatNoGeneric` on 19 Feb 2020

Conversation 2

Commits 1

Checks 84

Files changed 5

ts2do commented on 19 Feb 2020

Contributor ...

- Make the fast path for Number.FormatXX & Number.TryFormatXX inlineable
- Make parameterless ToString in SByte, Int16, Int32, and Int64 directly invoke the default formatting method

This includes the changes from [#32413](#) without the use of generics.

@jkotas: Can [#32413](#) be used as the 'other' PR or should I remove the optimizations and open another?



1



1

[Benchmark]

```
public string Roundtrip()
{
    byte[] bytes = Encoding.UTF8.GetBytes("this is a test");
    return Encoding.UTF8.GetString(bytes);
}
```

Method	Runtime	Mean	Ratio	Allocated
Roundtrip	.NET FW 4.8	113.69 ns	1.00	96 B
Roundtrip	.NET Core 3.1	49.76 ns	0.44	96 B
Roundtrip	.NET 5.0	36.70 ns	0.32	96 B

Improve Encoding.UTF8.GetString / GetChars performance for small inputs

#27268

Open with ▾

Merged

GrabYourPitchfo... merged 3 commits into [dotnet:master](#) from [GrabYourPitchforks:utf8_perf](#) on 9 Nov 2019

Conversation 18

-o Commits 3

Checks 52

Files changed 3

+120 -10



GrabYourPitchforks commented on 17 Oct 2019

Member



...

This improves the performance of `Encoding.UTF8.GetString(byte[])` : string and `Encoding.UTF8.GetBytes(string)` : byte[] by building on the existing JIT devirtualization logic and taking advantage of the case that most inputs to these functions are likely to be small (32 elements or fewer). For small inputs such as these, we already know that the maximum input size fits nicely into a stackalloc, so we can avoid the counting step and move straight to transcoding + the final memcpy.

Method	Toolchain	Text	Mean	Error	StdDev	Median	Ratio	Rate
GetString_FromByteArray	master		1.098 ns	0.0619 ns	0.0549 ns	1.075 ns	1.00	
GetString_FromByteArray	proto		8.682 ns	0.2368 ns	0.1978 ns	8.664 ns	7.91	
GetByteArray_FromString	master		21.109 ns	0.5399 ns	1.4032 ns	20.549 ns	1.00	
GetByteArray_FromString	proto		8.081 ns	0.1581 ns	0.1401 ns	8.025 ns	0.36	

Reviewers



EgorBo



gfoidl



tarekgh



jkotas



Assignees

No one assigned

Labels

area-System.Runtime

Projects

None yet

```

private readonly string _input = new HttpClient().GetStringAsync("http://www.gutenbe
private Regex _regex;

[Params(false, true)]
public bool Compiled { get; set; }

[GlobalSetup]
public void Setup() => _regex = new Regex(@"^.*\blove\b.*$", RegexOptions.Multiline);

[Benchmark]
public int Count() => _regex.Matches(_input).Count;

```

Method	Runtime	Compiled	Mean	Ratio
Count	.NET FW 4.8	False	26.207 ms	1.00
Count	.NET Core 3.1	False	21.106 ms	0.80
Count	.NET 5.0	False	4.065 ms	0.16
Count	.NET FW 4.8	True	16.944 ms	1.00
Count	.NET Core 3.1	True	15.287 ms	0.90
Count	.NET 5.0	True	2.172 ms	0.13

Optimize newline handling for RegexOptions.Multiline #34566

Merged

stephentoub merged 2 commits into `dotnet:master` from `stephentoub:bolffc`  on 9 Apr 2020

Conversation 9

Commits 2

Checks 126

Files changed 8



stephentoub commented on 5 Apr 2020 · edited

Member

...

We previously didn't do any special handling of beginning-of-line anchors (`^` when `RegexOptions.Multiline` is specified). This PR adds special handling for the anchor so that `FindFirstChar` will jump to the next newline as part of its processing.

As part of this, I also cleaned up some of the anchor handling code. The `RegexPrefixAnalyzer` only ever returns a single anchor, but the rest of the code was written such that it was expecting multiple anchors.

Example:

Finding all lines in Romeo and Juliet that contain the word "love":

Reduce unnecessary Regex match attempts for expressions beginning with atomic loops #35824

Open with ▾

Merged

danmosemsft merged 2 commits into `dotnet:master` from `stephentoub:atomicfast` on 5 May 2020

Conversation 7

Commits 2

Checks 116

Files changed 9

+137 -6



stephentoub commented on 5 May 2020 · edited

Member

...

Optimize Regex expressions that begin with atomic unbounded loops (either as written by the dev or because the system detected the loop could be made atomic and did it implicitly) by updating the starting position in the scan loop for the next iteration to be where the loop ended rather than where it started.

Running the examples from <https://github.com/mariomka/regex-benchmark/blob/652d55810691ad88e1c2292a2646d301d3928903/csharp/Benchmark.cs#L20-L26>:

Reviewers

danmosemsft



Assignees

No one assigned

```
private Regex _email = new Regex(@"[\w\.-]+@[^\w\.-]+\.[\w\.-]+", RegexOptions.Compiled);
private Regex _uri = new Regex(@"[\w]+://[^/\s?#]+[^/\s?#]+(?:\?:[^/\s?#]*)(?:#[^/\s]*)(?:[^\s])");
private Regex _ip = new Regex(@"(?:25[0-5]|2[0-4][0-9]|0[1-9][0-9])\.\){3}(?:25[0-5]|2[0-4][0-9]|0[1-9][0-9])\.\){1}(?:25[0-5]|2[0-4][0-9]|0[1-9][0-9])\.\){1}(?:25[0-5]|2[0-4][0-9]|0[1-9][0-9])");

private string _input = new HttpClient().GetStringAsync("https://raw.githubusercontent.com/");

[Benchmark] public int Email() => _email.Matches(_input).Count;
[Benchmark] public int Uri() => _uri.Matches(_input).Count;
[Benchmark] public int IP() => _ip.Matches(_input).Count;
```

Method	Runtime	Mean	Ratio
Email	.NET FW 4.8	1,036.729 ms	1.00
Email	.NET Core 3.1	930.238 ms	0.90
Email	.NET 5.0	50.911 ms	0.05
Uri	.NET FW 4.8	870.114 ms	1.00
Uri	.NET Core 3.1	759.079 ms	0.87
Uri	.NET 5.0	50.022 ms	0.06
IP	.NET FW 4.8	75.718 ms	1.00
IP	.NET Core 3.1	61.818 ms	0.82
IP	.NET 5.0	6.837 ms	0.09

Коллекции

Improve performance of iterating with ForEach over ImmutableList #1183

Open with ▾

Merged

jkotas merged 2 commits into `dotnet:master` from `hnrbaggio:iterateforeach_immutablearray` on 27 Dec 2019

Conversation 4

Commits 2

Checks 71

Files changed 1

+2 -0

Changes from all commits ▾

File filter... ▾

Jump to... ▾



0 / 1 files viewed



Review changes

2 ...System.Collections.Immutable/src/System/Collections/Immutable/ImmutableArray_1.Minimal.cs

Viewed

...

@@ -8,6 +8,7 @@

```
8     8     using System.Diagnostics.Contracts;
9     9     using System.Globalization;
10    10    using System.Linq;
11    11    + using System.Runtime.CompilerServices;
12    12    using System.Runtime.Versioning;
13    13
14    14    namespace System.Collections.Immutable
```

@@ -283,6 +284,7 @@ public ImmutableList<T>.Builder ToBuilder()

```
283   284        /// </summary>
284   285        /// <returns>An enumerator.</returns>
285   286        [Pure]
286   287    +    [MethodImpl(MethodImplOptions.AggressiveInlining)]
287   288        public Enumerator GetEnumerator()
288   289        {
289   290            var self = this;
```

```
private ImmutableList<int> _array = ImmutableList.Create(Enumerable.Range(0, 100_000).ToList());  
  
[Benchmark]  
public int Sum()  
{  
    int sum = 0;  
  
    foreach (int i in _array)  
        sum += i;  
  
    return sum;  
}
```

Method	Runtime	Mean	Ratio
Sum	.NET FW 4.8	187.60 us	1.00
Sum	.NET Core 3.1	187.32 us	1.00
Sum	.NET 5.0	46.59 us	0.25

```

private ImmutableList<int> _list = ImmutableList.Create(Enumerable.Range(0, 1_000).T

[Benchmark]
public int Sum()
{
    int sum = 0;

    for (int i = 0; i < 1_000; i++)
        if (_list.Contains(i))
            sum += i;

    return sum;
}

```

Method	Runtime	Mean	Ratio
Sum	.NET FW 4.8	22.259 ms	1.00
Sum	.NET Core 3.1	22.872 ms	1.03
Sum	.NET 5.0	2.066 ms	0.09

```

[GlobalSetup]
public void Setup()
{
    var r = new Random(42);
    _array = Enumerable.Range(0, 1_000).Select(_ => r.Next()).ToArray();
}

private int[] _array;

[Benchmark]
public void Sort()
{
    foreach (int i in _array.OrderBy(i => i)) { }
}

```

Method	Runtime	Mean	Ratio
Sort	.NET FW 4.8	100.78 us	1.00
Sort	.NET Core 3.1	101.03 us	1.00
Sort	.NET 5.0	85.46 us	0.85

Avoid mod operator when fast alternative available #27299

Merged

jkotas merged 12 commits into `dotnet:master` from `benaadams:dict-divisor`  on 26 Oct 2019

Conversation 86

Commits 12

Checks 52

Files changed 3



benaadams commented on 19 Oct 2019 • edited

Member

...

Use fastmod when 128bit multiply is available. (by @lemire <https://lemire.me/blog/2019/02/08/faster-remainders-when-the-divisor-is-a-constant-beating-compilers-and-libdivide/>)

Api for review dotnet/corefx#41822 to enable 128bit multiply and make it more efficient.

Other half of PR from [#27149](#) (first half was [#27195](#))

```
int i = buckets[hashCode % (uint)buckets.Length];
int i = GetBucket(hashCode);
```

```
int i = buckets[hashCode % (uint)buckets.Length];
int i = GetBucket(hashCode);
```

```
[MethodImpl(MethodImplOptions.AggressiveInlining)]
private ref int GetBucket(uint hashCode)
{
    int[] buckets = _buckets!;
#if BIT64
    return ref buckets[HashHelpers.FastMod(hashCode, (uint)buckets.Length, _fastModMultiplier)];
#else
    return ref buckets[hashCode % (uint)buckets.Length];
#endif
}
```

Daniel Lemire's blog

Daniel Lemire is a computer science professor at the University of Quebec (TELUQ) in Montreal. His research is focused on software performance and data engineering. He is a technooptimist.

[My home page](#)

[My papers](#)

[My software](#)

SUBSCRIBE

You can [subscribe to this blog by email](#).

WHERE TO FIND ME?

I am on Twitter and GitHub:



Faster remainders when the divisor is a constant: beating compilers and libdivide

Not all instructions on modern processors cost the same. Additions and subtractions are cheaper than multiplications which are themselves cheaper than divisions. For this reason, compilers frequently replace division instructions by multiplications. Roughly speaking, it works in this manner. Suppose that you want to divide a variable n by a constant d . You have that $n/d = n * (2^N/d) / (2^N)$. The division by a power of two ($/ (2^N)$) can be implemented as a right shift if we are working with unsigned integers, which compiles to single instruction: that is possible because the underlying hardware uses a base 2. Thus if $2^N/d$ has been pre-computed, you can compute the division n/d as a multiplication and a shift. Of course, if d is not a power of two, $2^N/d$ cannot be represented as an integer. Yet for N large enough^{footnote}, we can approximate $2^N/d$ by an integer and have the exact computation of the remainder for all possible n within a range. I believe that all optimizing C/C++ compilers know how to pull this trick and it is generally beneficial irrespective of the processor's architecture.

The idea is not novel and goes back to at least 1973 (Jacobsohn). However, engineering matters because computer registers have finite number of bits, and multiplications can overflow. I believe that, historically, this

```

private Dictionary<int, int> _dictionary = Enumerable.Range(0, 10_000).ToDictionary(
    [Benchmark]
public int Sum()
{
    Dictionary<int, int> dictionary = _dictionary;
    int sum = 0;

    for (int i = 0; i < 10_000; i++)
        if (dictionary.TryGetValue(i, out int value))
            sum += value;

    return sum;
}

```

Method	Runtime	Mean	Ratio
Sum	.NET FW 4.8	77.45 us	1.00
Sum	.NET Core 3.1	67.35 us	0.87
Sum	.NET 5.0	44.10 us	0.57

Rewrite HashSet<T>'s implementation based on Dictionary<T>'s #37180

Open with ▾

Merged stephentoub merged 9 commits into [dotnet:master](#) from [stephentoub:hashsetperf](#) on 1 Jun 2020

Conversation 52

Commits 9

Checks 186

Files changed 15

+2,003 -2,399



stephentoub commented on 30 May 2020

Member ...

Fixes [#37111](#)

Contributes to [#1989](#)

This moves HashSet into corelib, and then effectively deletes HashSet's data structure and replaces it with the one used by Dictionary, then updated for the differences (e.g. just a value rather than a key and a value). HashSet used to have basically the same implementation, but Dictionary has evolved significantly and HashSet hasn't; this brings them to basic parity on implementation.

Based on perf tests, I veered away from Dictionary's implementation in a few places (e.g. a goto-based implementation in the core find method led to a significant regression for Int32-based Contains operations), and we should follow-up to understand whether Dictionary should be changed as well, or why there's a difference between the two. @benaadams, if you have some time, it'd probably worth looking at this again; maybe you'll get different numbers than I did.

Functionally, bringing over Dictionary's implementation yields a few notable changes, namely that Remove and Clear no longer invalidate enumerations. The tests have been updated accordingly.

With HashSet now in corelib, I also updated two Dictionary uses I found in corelib that were using Dictionary as a set and switched them to use HashSet.

Running the dotnet/performance perf tests:

```
dotnet run -c Release -f net5.0 --filter System.Collections.*.HashSet --corerun d:\coreclrtest\master\corerun.exe
```

Reviewers

danmosemsft ✓

jkotas

AntonLapounov

Assignees

No one assigned

Labels

[area-System.Collections](#)

[tenant-performance](#)

Projects

None yet

Milestone

5.0.0

New serializer converter model for objects and collections #2259

Open with ▾

Merged

steveharter merged 3 commits into `dotnet:master` from `steveharter:NewConverterModel` on 12 Feb 2020

Conversation 507

Commits 3

Checks 66

Files changed 108

+6,049 -4,906



steveharter commented on 28 Jan 2020 · edited by richlander

Member

...

As part of #1562 this PR refactors the converter model for objects and collection.

No public APIs have changed with this PR - changes were made to be internal for now.

Short-term this PR improves perf and maintainability.

Longer-term a subsequent PR will be created for the proposed public API changes which are expected to add capabilities as explained in the above link (mostly supporting code-gen for AOT and adding public APIs for extensibility and new converter capabilities).

Reviewers

layomia



Jozkee



safern



ahsonkhan



Assignees

steveharter

```

private MemoryStream _stream = new MemoryStream();
private DateTime[] _array = Enumerable.Range(0, 1000).Select(_ => DateTime.UtcNow).T

[Benchmark]
public Task LargeArray()
{
    _stream.Position = 0;
    return JsonSerializer.SerializeAsync(_stream, _array);
}

```

Method	Runtime	Mean	Ratio	Allocated
LargeArray	.NET FW 4.8	262.06 us	1.00	24256 B
LargeArray	.NET Core 3.1	191.34 us	0.73	24184 B
LargeArray	.NET 5.0	69.40 us	0.26	152 B

A solid red diagonal bar runs from the top-left corner to the bottom-left corner of the image.

Сеть

[Benchmark]

```
public Uri Ctor() => new Uri("https://github.com/dotnet/runtime/pull/36915");
```

Method	Runtime	Mean	Ratio	Allocated
Ctor	.NET FW 4.8	443.2 ns	1.00	225 B
Ctor	.NET Core 3.1	192.3 ns	0.43	72 B
Ctor	.NET 5.0	129.9 ns	0.29	56 B

```
private Uri _uri = new Uri("http://github.com/dotnet/runtime");
```

[Benchmark]

```
public string PathAndQuery() => _uri.PathAndQuery;
```

Method	Runtime	Mean	Ratio	Allocated
PathAndQuery	.NET FW 4.8	17.936 ns	1.00	56 B
PathAndQuery	.NET Core 3.1	30.891 ns	1.72	56 B
PathAndQuery	.NET 5.0	2.854 ns	0.16	—

```
using Microsoft.AspNetCore.Builder;
using Microsoft.AspNetCore.Hosting;
using Microsoft.AspNetCore.Http;
using Microsoft.Extensions.Hosting;

public class Program
{
    public static void Main(string[] args) =>
        Host.CreateDefaultBuilder(args).ConfigureWebHostDefaults(b => b.UseStartup<Startup>());
}

public class Startup
{
    public void Configure(IApplicationBuilder app, IWebHostEnvironment env)
    {
        app.UseRouting();
        app.UseEndpoints(endpoints =>
        {
            endpoints.MapGet("/", context => context.Response.WriteAsync("Hello"));
            endpoints.MapPost("/", context => context.Response.WriteAsync("Hello"));
        });
    }
}
```

```

private HttpResponseMessage _client = new HttpResponseMessage(new SocketsHttpHandler())
private HttpRequestMessage _get = new HttpRequestMessage(HttpMethod.Get, new Uri("ht
private HttpRequestMessage _post = new HttpRequestMessage(HttpMethod.Post, new Uri("

[Benchmark] public Task Get() => MakeRequest(_get);

[Benchmark] public Task Post() => MakeRequest(_post);

private Task MakeRequest(HttpRequestMessage request) => Task.WhenAll(Enumerable.Rang
{
    for (int i = 0; i < 500; i++)
    {
        using (HttpResponseMessage r = await _client.SendAsync(request, default))
        using (Stream s = await r.Content.ReadAsStreamAsync())
            await s.CopyToAsync(Stream.Null);
    }
});

```

Method	Runtime	Mean	Ratio	Allocated
Get	.NET Core 3.1	1,267.4 ms	1.00	122.76 MB
Get	.NET 5.0	681.7 ms	0.54	74.01 MB
Post	.NET Core 3.1	1,464.7 ms	1.00	280.51 MB
Post	.NET 5.0	735.6 ms	0.50	132.52 MB

Потоки и асинхроннос ть

[Benchmark]

```
public async Task ValueTaskCost()
```

```
{  
    for (int i = 0; i < 1_000; i++)  
        await YieldOnce();  
}
```

```
private static async ValueTask YieldOnce() => await Task.Yield();
```

Method	Runtime	Mean	Ratio	Allocated
ValueTaskCost	.NET FW 4.8	1,635.6 us	1.00	294010 B
ValueTaskCost	.NET Core 3.1	842.7 us	0.51	120184 B
ValueTaskCost	.NET 5.0	812.3 us	0.50	186 B

Make "async ValueTask/ValueTask<T>" methods ammortized allocation-free

#26310

Open with ▾

Merged

stephentoub merged 1 commit into [dotnet:master](#) from [stephentoub:asyncvaluetaskallocs](#) on 24 Oct 2019

Conversation 84

Commits 1

Checks 52

Files changed 7

+654 -139



stephentoub commented on 22 Aug 2019 • edited

Member



...

Today `async ValueTask/ValueTask<T>` methods use builders that special-case the synchronously completing case (to just return a `default(ValueTask)` or new `ValueTask<T>(result)`) but defer to the equivalent of `async Task/Task<T>` for when the operation completes asynchronously. This, however, doesn't take advantage of the fact that value tasks can wrap arbitrary `IValueTaskSource/IValueTaskSource<T>` implementations.

I've had this sitting on the shelf for a while, but finally cleaned it up. The first three commits here are just moving around existing code. The last two commits are the meat of this change. This changes `AsyncValueTaskMethodBuilder` and `AsyncValueTaskMethodBuilder<T>` to use pooled `IValueTaskSource/IValueTaskSource<T>` instances, such that calls to an `async ValueTask/ValueTask<T>` method incur 0 allocations as long as there's a cached object available.

I've marked this as a draft and work-in-progress for a few reasons:

1. There's a breaking change here, in that while we say/document that `ValueTask/ValueTask<T>`s are more limited in what they can be used for, nothing in the implementation actually stops a `ValueTask` that was wrapping a `Task` from being used as permissively as `Task`, e.g. if you were to await such a `ValueTask` multiple times, it would happen to work, even though we say "never do that". This change means that anyone who was relying on such undocumented behaviors will now be broken. I think this is a reasonable thing to do in a major release, but I also want feedback and a lot of runway on it. There's the potential to make it opt-in (or opt-out) as well, but that will also non-trivially complicate the implementation.
2. Policy around pooling. This is always a difficult thing to tune. Right now I've chosen a policy that limits the number of pooled objects per state machine type to an arbitrary multiple of the processor count, and that tries to strike a balance between contention and garbage by using a spin lock and if there's any contention while trying to get or return a pooled object, the cache is ignored. We will need to think hard about what policy to use here. It's also possible it could be tuned per state machine, e.g. by having an attribute that's looked up via reflection when creating the cache for a state machine, but that'll add a lot of first-access overhead to any `async ValueTask/ValueTask<T>` method.
3. Perf validation.

cc: @kouvel, @tarekgh, @benaadams

30

14

github.com/dotnet/coreclr/pull/26310

Reviewers

davidfowl

benaadams



Assignees

No one assigned

Labels

area-System.Threading

Projects

None yet

Milestone

5.0

Linked issues

Successfully merging this pull request may close these issues.

None yet

Notifications

Customize

Subscribe

You're not receiving notifications from this thread.

Async ValueTask Pooling in .NET 5



Stephen

March 16th, 2020

The `async/await` feature in C# has revolutionized how developers targeting .NET write asynchronous code. Sprinkle some `async` and `await` around, change some return types to be tasks, and badda bing badda boom, you've got an asynchronous implementation. In theory.

In practice, obviously I've exaggerated the ease with which a codebase can be made fully asynchronous, and as with many a software development task, the devil is often in the details. One such "devil" that performance-minded .NET developers are likely familiar with is the state machine object that enables an `async` method to perform its magic.

Improve ContinueWith perf with NotOn* options #35575

Open with ▾

Merged stephentoub merged 5 commits into `dotnet:master` from `stephentoub:cwperf` on 2 May 2020

Conversation 2

-o- Commits 5

Checks 129

Files changed 2

+87 -95



stephentoub commented on 28 Apr 2020

Member

...

This came about while looking at allocations and CPU costs in some HttpClient code. When ContinueWith is used with TaskContinuationOptions.NotOn* options, when the antecedent task completes we compare the state of that antecedent task against the NotOn* options: if the options permit it, the continuation is queued/invoked, and if they don't, the continuation is canceled. That cancellation then ends up being common in cases where a ContinueWith is used, for example, to log exceptions that result from a faulted task, e.g.

```
private static void Log(Task task) =>
    task.ContinueWith(t => Log(t.Exception),
        CancellationToken.None,
        TaskContinuationOptions.OnlyOnFaulted | TaskContinuationOptions.ExecuteSynchronously,
        TaskScheduler.Default);
```

We can handle that cancellation much more efficiently than we are today. Today that's resulting in us expanding the ContingentProperties inside of the task, in order to store that cancellation was requested internally, but that's not actually necessary. It's also resulting in us doing several atomic transitions via interlocks, but that's only necessary if the task could have transitioned in any way, which is only possible if a cancelable token is provided.

This PR removes that overhead. It also shrinks the size of the object created by ContinueWith by a field, seals and renames it, and removes some dead code from related code paths.

Reviewers

tarekgh ✓

Assignees

No one assigned

Labels

area-System.Threading.Tasks
tenet-performance

Projects

None yet

Milestone

5.0.0

```

const int Iters = 1_000_000;

private AsyncTaskMethodBuilder[] tasks = new AsyncTaskMethodBuilder[Iter];

[IterationSetup]
public void Setup()
{
    Array.Clear(tasks, 0, tasks.Length);
    for (int i = 0; i < tasks.Length; i++)
        _ = tasks[i].Task;
}

[Benchmark(OperationsPerInvoke = Iter)]
public void Cancel()
{
    for (int i = 0; i < tasks.Length; i++)
    {
        tasks[i].Task.ContinueWith(_ => { }, CancellationToken.None, TaskContinuationOptions
            tasks[i].SetResult();
    }
}

```

Method	Runtime	Mean	Ratio	Allocated
Cancel	.NET FW 4.8	239.2 ns	1.00	193 B
Cancel	.NET Core 3.1	140.3 ns	0.59	192 B
Cancel	.NET 5.0	106.4 ns	0.44	112 B



JIT

```

[Benchmark]
public int Zeroing()
{
    ReadOnlySpan<char> s1 = "hello world";
    ReadOnlySpan<char> s2 = Nop(s1);
    ReadOnlySpan<char> s3 = Nop(s2);
    ReadOnlySpan<char> s4 = Nop(s3);
    ReadOnlySpan<char> s5 = Nop(s4);
    ReadOnlySpan<char> s6 = Nop(s5);
    ReadOnlySpan<char> s7 = Nop(s6);
    ReadOnlySpan<char> s8 = Nop(s7);
    ReadOnlySpan<char> s9 = Nop(s8);
    ReadOnlySpan<char> s10 = Nop(s9);
    return s1.Length + s2.Length + s3.Length + s4.Length + s5.Length + s6.Length + s
}

[MethodImpl(MethodImplOptions.NoInlining)]
private static ReadOnlySpan<char> Nop(ReadOnlySpan<char> span) => default;

```

Method	Runtime	Mean	Ratio
Zeroing	.NET FW 4.8	22.85 ns	1.00
Zeroing	.NET Core 3.1	18.60 ns	0.81
Zeroing	.NET 5.0	15.07 ns	0.66

Use xmm for stack prolog zeroing rather than rep stos #32538

Merged AndyAyersMS merged 17 commits into `dotnet:master` from `benaadams:zero-init-xmm` on 5 Mar 2020

Conversation 217 · Commits 17 · Checks 149 · Files changed 8

benaadams commented on 19 Feb 2020 · edited

Member ...

Use `xmm` registers to clear prolog blocks rather than `rep stosd`.

`rep stos` does have a shallower growth rate so it will eventually overtake; however unlikely in the size range for the stack clearance:

488 bytes zero'd same cost as 40 bytes previously

Also its variability based on small changes in stack changes is problematic as just adding an extra variable to a method (with no other changes) could add or remove 9ns from its running time. Whereas the xmm clear in this PR is much smoother and more expected cost.

github.com/dotnet/runtime/pull/32538

Optimization to remove redundant zero initializations. #36918

Merged

erzenfeld merged 2 commits into `dotnet:master` from `erzenfeld:ZeroInitDiff` on 27 May 2020



Conversation 31

Commits 2

Checks 177

Files changed 10



erzenfeld commented on 23 May 2020

Member

...

This change adds a phase that iterates over basic blocks starting with the first basic block until there is no unique basic block successor or until it detects a loop. It keeps track of local nodes it encounters. When it gets to an assignment to a local variable or a local field, it checks whether the assignment is the first reference to the local (or to the parent of the local field), and, if so, it may do one of two optimizations:

1. If the local is untracked, the rhs of the assignment is 0, and the local is guaranteed to be fully initialized in the prolog, the explicit zero initialization is removed.
2. If the assignment is to a local (and not a field) and either the local has no gc pointers or there are no gc-safe points between the prolog and the assignment, it marks the local with `lvHasExplicitInit` which tells the codegen not to insert zero initialization for this local in the prolog.

This addresses one of the examples in [#2325](#) and 5 examples in [#1007](#).



13



5

```
public static char Get(string s, int i) => s[i];
```

```
; Program.Get(System.String, Int32)
    sub      rsp,28
    cmp      edx,[rcx+8]
    jae      short M01_L00
    movsxd  rax,edx
    movzx   eax,word ptr [rcx+rax*2+0C]
    add      rsp,28
    ret
M01_L00:
    call    CORINFO_HELP_RNGCHKFAIL
    int     3
; Total bytes of code 28
```

```
public static char Get(string s, int i) => s[i];
```

```
; Program.Get(System.String, Int32)
    sub      rsp,28
    cmp      edx,[rcx+8]
    jae      short M01_L00
    movsxd  rax,edx
    movzx   eax,word ptr [rcx+rax*2+0C]
    add      rsp,28
    ret
M01_L00:
    call    CORINFO_HELP_RNGCHKFAIL
    int     3
; Total bytes of code 28
```

```
private (int i, int j) _value;
```

[Benchmark]

```
public int NullCheck() => _value.j++;
```

```
; Program.NullCheck()
```

```
    nop      dword ptr [rax+rax]
    cmp      [rcx],ecx
    add      rcx,8
    add      rcx,4
    mov      eax,[rcx]
    lea      edx,[rax+1]
    mov      [rcx],edx
    ret
```

```
; Total bytes of code 23
```

```
private (int i, int j) _value;
```

[Benchmark]

```
public int NullCheck() => _value.j++;
```

```
; Program.NullCheck()
```

```
    nop      dword ptr [rax+rax]
    cmp      [rcx],ecx
    add      rcx,8
    add      rcx,4
    mov      eax,[rcx]
    lea      edx,[rax+1]
    mov      [rcx],edx
    ret
```

```
; Total bytes of code 23
```

```
private (int i, int j) _value;
```

[Benchmark]

```
public int NullCheck() => _value.j++;
```

```
; Program.NullCheck()
```

```
    nop      dword ptr [rax+rax]
    cmp      [rcx],ecx
    add      rcx,8
    add      rcx,4
    mov      eax,[rcx]
    lea      edx,[rax+1]
    mov      [rcx],edx
    ret
```

```
; Total bytes of code 23
```

```
private (int i, int j) _value;
```

[Benchmark]

```
public int NullCheck() => _value.j++;
```

```
; Program.NullCheck()
```

```
    add    rcx, 0C  
    mov    eax, [rcx]  
    lea    edx, [rax+1]  
    mov    [rcx], edx  
    ret
```

```
; Total bytes of code 12
```

```
private int[] _array = Enumerable.Range(0, 1000).ToArray();

[Benchmark]
public bool IsSorted() => IsSorted(_array);

private static bool IsSorted(ReadOnlySpan<int> span)
{
    for (int i = 0; i < span.Length - 1; i++)
        if (span[i] > span[i + 1])
            return false;

    return true;
}
```

```

private int[] _array = Enumerable.Range(0, 1000).ToArray();

[Benchmark]
public bool IsSorted() => IsSorted(_array);

private static bool IsSorted(ReadOnlySpan<int> span)
{
    for (int i = 0; i < span.Length - 1; i++)
        if (span[i] > span[i + 1])
            return false;

    return true;
}

```

Method	Runtime	Mean	Ratio	Code Size
IsSorted	.NET FW 4.8	1,083.8 ns	1.00	236 B
IsSorted	.NET Core 3.1	581.2 ns	0.54	136 B
IsSorted	.NET 5.0	463.0 ns	0.43	105 B

Рома Неволин



nevoroman@gmail.com



t.me/nevoroman