

Span, Memory and Pipelines, the APIs you always missed



Raffaele Rialdi - Senior Software Architect



@raffaeler
raffaeler@vevy.com

Who am I?



- Raffaele Rialdi, Senior Software Architect in Vevy Europe – Italy
 - @raffaeler also known as "Raf"
- Consultant in many industries
 - Manufacturing, racing, healthcare, financial, ...
- Speaker and Trainer around the globe (development and security)
 - Italy, Romania, Bulgaria, Russia (Moscow, St Petersburg and Novosibirsk), USA, ...
- And proud member of the great Microsoft MVP family since 2003



Agenda

- Using value type by reference is the key
- Our new best friends: `Span<T>` and `Memory<T>`
- Going unsafe
- The new memory allocation primitives
- Pipelines: a better way to manage streams of data
- Realtime processing with `Span`, `Memory` and `Pipelines`

A long dated problem

- Value-based vs Reference-based languages
- .NET is value-based but splits the type system in value and reference types

	Reference Types	Value Types
Allocation	heap (GC involved) ⚠	stack (no GC)
What is copied	just the reference	the whole data ⚠

- C# expanded the ability to work with references on value types
 - Started with C#7 and continued in C#8

Less GC, more performance, still safe

- C# 7.x widened the reference paradigm to avoid copies
 - the "in" modifier, meaning "readonly ref"
 - using "ref" when returning values
 - declaring local "ref" or "readonly ref" local variables
- The new «ref struct» and «readonly ref struct» ensure at compile time that instances will only live on the stack (no GC involved)
- Using ref is like viewing memory without owning it
 - You can both read and write it, provided it is not readonly

public readonly **ref** partial **struct** Span<T>

Span<T> and ReadOnlySpan<T>

- They are both "ref readonly struct"
 - The compiler ensure they only live on the stack, not hitting the GC at all
- It is a "view" over a contiguous region of memory
 - Every change on the view is effectively made on the memory being viewed

```
Span<byte> span = new byte[]  
    { 0, 2, 4, 6, 8, 10, 12, 14, 16 }.AsSpan();  
int total = 0;  
foreach (byte item in span.Slice(3, 5))  
    total += item;  
Debug.Assert(total == 50);
```

```
string hello = "Hello, world!";  
ReadOnlySpan<char> span1 = hello;  
ReadOnlySpan<char> span2 = span1.Slice(7, 5);  
  
Debug.Assert(span2.ToString() == "world");  
Debug.Assert(span2 != "world");
```

- Designed to easily wrap any array

0	1	2	3	4	5	6	7	8	9	A	B	C	D	----
	0		1		2		3		4		5		6	----
00	48	00	65	00	6C	00	6C	00	6F	00	20	00	2C	----
	H		e		l		l		o				,	----

Example: A no-GC version of the string.Trim()

```
string test = "    Hello, World! ";  
Trim(test).ToArray()
```

Span<T>

- is allocated **on the stack**
- cannot be stored as a class member
- does not involve any heap allocation
- does not impact on **GC**
- is a view on **managed** or **native** memory

An immutable view over a string

```
ReadOnlySpan<char> Trim(ReadOnlySpan<char> source)  
{  
    if (source.IsEmpty) return source;  
    int start = 0, end = source.Length - 1;  
    char startChar = source[start]  
    char endChar = source[end];  
    while ((start < end) &&  
           (startChar == ' ' || endChar == ' '))  
    {  
        if (startChar == ' ') start++;  
        if (endChar == ' ') end--;  
        startChar = source[start];  
        endChar = source[end];  
    }  
    return source.Slice(start, end - start + 1);  
}
```

A new immutable view over a string


Span<T> limitations


- ref struct are allowed only in ref structs
- can't declare ref struct in async methods, but ... look at the example!
- As it is a ref struct, can't survive the stack unwind in local functions


```
private async Task SomeAsyncFunc()
{
    var memory = new Memory<byte>(new byte[100]);
    await Task.Delay(1);

    // Not allowed in async methods
    //var span = memory.Span;
    //ref var a = ref MyLocalFunc1();

    MyLocalFunction() = 99;
    ref byte MyLocalFunction() => ref memory.Span[1];
}
```



 `public class SomeClass`
`{`
 `Span<byte> span;`
`}`

 `public ref struct SomeStruct`
`{`
 `Span<byte> span;`
`}`

Memory<T>

```
[DebuggerTypeProxy(typeof(MemoryDebugView<>))]  
[DebuggerDisplay("{ToString(),raw}")]  
public readonly struct Memory<T>
```

- Wraps a contiguous block of memory by holding a reference to it
 - It is **not** a ref struct and can survive stack unwind
 - The ***Span property*** expose a "view" of the memory hold by Memory<T>
 - Span<T> can't be converted in Memory<T> (a **copy** is needed)
- Rich extension methods provided in the box
 - AsSpan, AsMemory, BinarySearch, IndexOf, LastIndexOf, ...

```
var m1 = new Memory<byte>();  
Debug.Assert(m1.IsEmpty);  
  
ReadOnlyMemory<char> memStr =  
    "Hello, world".AsMemory();
```

```
var blob = new byte[100];  
  
var m2 = new Memory<byte>(blob);  
var m3 = new Memory<byte>(blob, start: 10, length: 5);  
var m4 = blob.AsMemory();  
var m5 = blob.AsMemory(start:10);  
var m6 = blob.AsMemory(start:10, length:5);
```

Span<T> on strings benchmark

- Using Benchmark.NET to measure trimming " Hello, world "

Method	Loop	Mean	Error	StdDev	Gen 0/1k Op	Allocated Memory/Op
StringTrim	1000	23.25 us	0.4561 us	0.4684 us	13.3362	56000 B
SpanTrim	1000	16.44 us	0.3164 us	0.4001 us	-	-

```
[Benchmark]
public void StringTrim()
{
    for(int i=0; i<Loop; i++)
    {
        string res = Text.Trim();
    }
}
```

```
[Benchmark]
public void SpanTrim()
{
    ReadOnlySpan<char> span = Text;
    for (int i = 0; i < Loop; i++)
    {
        ReadOnlySpan<char> res = span.Trim();
    }
}
```

Span<T> and Unsafe

Span<T> and pointers

- Span<T> can be used on unsafe, classic pointers (byte *, ...)
 - Unsafe code is limited to construction, the rest is safe!
- Get some raw pointer

```
byte* ptr = _native.ReadUnsafe();
```

```
Span<byte> spanByte = _native.ReadUnsafe();
```

- Build a Span<byte>

```
Span<byte> spanByte = new Span<byte>(ptr, sizeof(WavHeader));
```

- Or a Span<WavHeader>

```
Span<WavHeader> spanHeader = new Span<WavHeader>(ptr, 1);
```

- We just *casted a managed struct to native memory allocation*

MemoryMarshal and Unsafe helper classes

- Casting a Span<byte> to a Span<T>

```
Span<WavHeader> spanHeader = MemoryMarshal.Cast<byte, WavHeader>(spanByte);
```



- Materializing an instance of T

```
WavHeader wavheader = MemoryMarshal.Read<WavHeader>(spanByte);
```

```
WavHeader wavheader = Unsafe.Read<WavHeader>(ptr);
```

- Avoid materialization getting just a reference to T

```
ref WavHeader refWavHeader = ref MemoryMarshal.GetReference<WavHeader>(spanHeader);
```

```
ref WavHeader refwavheader = ref Unsafe.AsRef<WavHeader>(ptr);
```

NetCore source code for AsRef



```
.maxstack 1  
ldarg.0  
ret
```

Introducing new memory management APIs

In the beginning ...

- ... we had classic allocation

```
byte[] blob = new byte[_size];
```

```
Memory<byte> memory = blob;
```

- Memory<byte> can encapsulate and manage the ownership
- Or we could allocate on the stack using unsafe

```
byte* ptr = stackalloc byte[size];
```

ArrayPool

- ArrayPool<T> allows renting and returning chunks of memory

```
byte[] blob = ArrayPool<byte>.Shared.Rent(size);
```

- Be careful on the returned size!

```
Debug.Assert(blob2a.Length >= _size);
```

- Be careful to **return** the rented buffer

```
ArrayPool<byte>.Shared.Return(blob, clearArray:false);
```

- Instead of the standard pool, we can create new ones

```
var mypool = ArrayPool<byte>.Create(  
    maxArrayLength: 1024,  
    maxArraysPerBucket: 10);
```


MemoryPool

- MemoryPool<T> is similar but supports the disposable pattern


```
using (IMemoryOwner<byte> blob = MemoryPool<byte>.Shared.Rent(size))
{
    Debug.Assert(blob.Memory.Length != size);
    // slicing is a good way to obtain the exact buffer size
    Memory<byte> memory = blob.Memory.Slice(0, size);
}
```

- You can create a custom pool by deriving MemoryPool<T>
 - Example on GitHub: ArrayMemoryPool<T>

Allocating on the stack

- With `Span<T>`, `stackalloc` does not require unsafe code anymore

```
Span<byte> blob = stackalloc byte[] { 0, 1, 2, 3, 4, 5 };
```




stackalloc initializer

- Avoid large buffers on the stack
 - C# stack default size is 1 MB
- This is the fastest possible allocation method for temporary buffers

ReadOnlySequence<T>

- ReadOnlySequence<T> is a linked list of memory segments/chunks
 - Each segment is made of contiguous memory
 - Segments are not (necessarily) contiguous in memory
 - Segments are exposed via Enumerator (do not implement IEnumerable<T>)
- Segments are anything deriving from ReadOnlySequenceSegment<T>
 - **ReadOnlySequenceSegment<T>** is abstract
 - There is no public concrete class available in corefx



The Pipeline API

Pipeline API

- Think to it as a modern Stream API
 - Conceptually mimes an (in-process) FIFO queue
 - Decouples readers from writers
 - Provides a built-in memory management for buffers
 - Leverages the power of:
 - `Span<T>`, `Memory<T>`, `MemoryPool<T>` and `ReadOnlySequence<T>`
 - Readers may decide to consume only a portion of the available buffer
- The content of the stream is always "bytes"

Writing a Pipe

- Strategy 1

- The Pipe uses a private Pool to rent segments of memory
- FlushAsync makes data available to the reader
- Memory is automatically returned as soon as the data is consumed

① *GetMemory strategy*

```
var pipe = new Pipe();

Memory<byte> mem = pipe.Writer.GetMemory(minSize)
int written = Encoding.UTF8.GetBytes("message",
                                     mem.Span);

pipe.Writer.Advance(written);
await pipe.Writer.FlushAsync();

pipe.Writer.Complete();
```

② *WriteAsync strategy*

```
var pipe = new Pipe();

Memory<byte> mem = Encoding.UTF8.GetBytes("message")
                                     .AsMemory();
var writeResult = await pipe.Writer.WriteAsync(mem);

pipe.Writer.Complete();
```

- Strategy 2

- The Pipe Writes an arbitrary blob of memory asynchronously

Reading a Pipe

- Strategy 1

- Usually used inside an infinite loop
- The async call is ended by completing the writer

- Strategy 2

- Used only when you need the current content (if any) without waiting

- The buffer is always a **ReadOnlySequence<byte>**

① *Asynchronous*

```
var result = await pipe.Reader.ReadAsync();  
var buffer = result.Buffer;  
if (result.IsCanceled || buffer.IsEmpty) {  
    // exit  
}
```

② *Synchronous*


```
if(!reader.TryRead(out ReadResult result) ||  
    result.IsCanceled || result.Buffer.IsEmpty) {  
    // exit  
}  
  
var buffer = result.Buffer;
```

Demo on Pipelines:
read a process stdout stream

To sum up

- .NET Core is finally mature and offers modern APIs
- Try moving the hot paths from the GC to the stack with `Span<T>`
- Slice buffers using `Span<T>`
- Use memory pools to minimize the GC cost on reusable buffers
- Evaluate replacing `System.IO.Stream` with Pipelines

Questions?



Questions @ booth 1

Thank you!